

Project

December 11, 2023

1 Project: F1 race prediction

1.1 1) Import necessary libraries

```
[1]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.preprocessing import OneHotEncoder
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split, GridSearchCV, KFold
from sklearn.linear_model import LinearRegression, Ridge, Lasso
from sklearn.metrics import r2_score, mean_squared_error
import xgboost as xgb
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor

import warnings
warnings.filterwarnings('ignore')
```

1.2 2) Helper Functions

```
[2]: # Convert timestamp to numeric value
def time_str_to_numeric(time_str):
    if isinstance(time_str, str) and time_str != "\\N":
        minutes, seconds = map(float, time_str.split(':'))
        return minutes * 60 + seconds
    else:
        return np.nan
```

```
[3]: # Map the race positions to the points
def attach_raceId_to_predictions(raceIds, y_pred):
    predictions_with_raceId = pd.DataFrame(raceIds)
    predictions_with_raceId['y_pred'] = y_pred
    predictions_with_raceId['new_y_pred'] = predictions_with_raceId.
    ↪groupby('raceId')['y_pred'].transform(lambda x: x.rank(ascending=False,
    ↪method='first'))
```

```

rank_to_points = {1: 25, 2: 18, 3: 15, 4: 12, 5: 10, 6: 8, 7: 6, 8: 4, 9: 2, 10: 1}

predictions_with_raceId['new_y_pred'] = predictions_with_raceId['new_y_pred'].map(rank_to_points).fillna(0)
new_y_pred = predictions_with_raceId['new_y_pred'].values

return new_y_pred

```

1.3 3) Data loading

1. All the csv files are in the data sub folder but we only need these 6 csv files

```

[4]: # Read the data from a csv files to dataframes
circuits = pd.read_csv("data/circuits.csv")
constructors = pd.read_csv("data/constructors.csv")
drivers = pd.read_csv("data/drivers.csv")
qualifying = pd.read_csv("data/qualifying.csv")
races = pd.read_csv("data/races.csv")
results = pd.read_csv("data/results.csv")

```

1.4 4) Data Pre-Processing

1.4.1 4.1) Raw data from multiple csv files

```

[5]: qualifying.head()

```

```

[5]:   qualifyId  raceId  driverId  constructorId  number  position  q1 \
0         1      18         1             1      22         1  1:26.572
1         2      18         9             2       4         2  1:26.103
2         3      18         5             1     23         3  1:25.664
3         4      18        13             6       2         4  1:25.994
4         5      18         2             2       3         5  1:25.960

      q2      q3
0  1:25.187  1:26.714
1  1:25.315  1:26.869
2  1:25.452  1:27.079
3  1:25.691  1:27.178
4  1:25.518  1:27.236

```

```

[6]: races.head()

```

```

[6]:   raceId  year  round  circuitId      name      date \
0      1  2009      1          1  Australian Grand Prix  2009-03-29
1      2  2009      2          2  Malaysian Grand Prix  2009-04-05
2      3  2009      3         17    Chinese Grand Prix  2009-04-19

```

3	4	2009	4	3	Bahrain Grand Prix	2009-04-26
4	5	2009	5	4	Spanish Grand Prix	2009-05-10

	time	url	fp1_date	\
0	06:00:00	http://en.wikipedia.org/wiki/2009_Australian_G...	\N	
1	09:00:00	http://en.wikipedia.org/wiki/2009_Malaysian_Gr...	\N	
2	07:00:00	http://en.wikipedia.org/wiki/2009_Chinese_Gran...	\N	
3	12:00:00	http://en.wikipedia.org/wiki/2009_Bahrain_Gran...	\N	
4	12:00:00	http://en.wikipedia.org/wiki/2009_Spanish_Gran...	\N	

	fp1_time	fp2_date	fp2_time	fp3_date	fp3_time	quali_date	quali_time	\
0	\N	\N	\N	\N	\N	\N	\N	
1	\N	\N	\N	\N	\N	\N	\N	
2	\N	\N	\N	\N	\N	\N	\N	
3	\N	\N	\N	\N	\N	\N	\N	
4	\N	\N	\N	\N	\N	\N	\N	

	sprint_date	sprint_time
0	\N	\N
1	\N	\N
2	\N	\N
3	\N	\N
4	\N	\N

```
[7]: results.head()
```

```
[7]:
```

	resultId	raceId	driverId	constructorId	number	grid	position	\
0	1	18	1	1	22	1	1	
1	2	18	2	2	3	5	2	
2	3	18	3	3	7	7	3	
3	4	18	4	4	5	11	4	
4	5	18	5	1	23	3	5	

	positionText	positionOrder	points	laps	time	milliseconds	\
0	1	1	10.0	58	1:34:50.616	5690616	
1	2	2	8.0	58	+5.478	5696094	
2	3	3	6.0	58	+8.163	5698779	
3	4	4	5.0	58	+17.181	5707797	
4	5	5	4.0	58	+18.014	5708630	

	fastestLap	rank	fastestLapTime	fastestLapSpeed	statusId
0	39	2	1:27.452	218.300	1
1	41	3	1:27.739	217.586	1
2	41	5	1:28.090	216.719	1
3	58	7	1:28.603	215.464	1
4	43	1	1:27.418	218.385	1

1.4.2 4.2) Data Selection, Handling missing data , Data Cleaning

```
[8]: # Remove unwanted columns
qualifying = qualifying.drop(['number'], axis = 1)
races = races[['raceId', 'year', 'round', 'circuitId']]
results = results.drop(['number', 'positionText', 'positionOrder', 'time', 'milliseconds', 'rank', 'statusId'], axis=1)

# Collecting races which happened on or after 2000 for our analysis
races = races[races['year'] >= 2000]
unique_race_ids = races['raceId'].unique()
results = results[results['raceId'].isin(unique_race_ids)]
qualifying = qualifying[qualifying['raceId'].isin(unique_race_ids)]

# Replacing the missing value of fastestLapTime with 100:0 (very large time value)
results['fastestLapTime'] = results['fastestLapTime'].replace({'\\N': '100:0'})

# Replacing the missing value of fastestLapSpeed with 0 (Least speed)
results['fastestLapSpeed'] = results['fastestLapSpeed'].replace({'\\N': '0'})

# Replacing the missing value of fastestLap with 0 (Indicating no lap)
results['fastestLap'] = results['fastestLap'].replace({'\\N': '0'})

# Replacing the missing value of position with 20 (last position)
results['position'] = results['position'].replace({'\\N': '20'})

# Converting the timestamp to numeric values
results['fastestLapTime'] = results['fastestLapTime'].apply(time_str_to_numeric)
qualifying['q1'] = qualifying['q1'].apply(time_str_to_numeric)
qualifying['q2'] = qualifying['q2'].apply(time_str_to_numeric)
qualifying['q3'] = qualifying['q3'].apply(time_str_to_numeric)

# Converting position, fastestLap, fastestLapSpeed from object to int
results['position'] = results['position'].astype(int)
results['fastestLap'] = results['fastestLap'].astype(int)
results['fastestLapSpeed'] = results['fastestLapSpeed'].astype(float)
```

1.4.3 4.3) Feature Engineering

1. Generate qualifying_time feature from q1, q2, q3
2. Generate features grid_pos1, grid_pos2, ..., grid_pos10 to capture past 10 grid positions of that driver
3. Generate features result_pos1, result_pos2, ..., result_pos10 to capture past 10 results positions of that driver
4. We believe that these features will give the models idea about the race car performance and driver skills

```
[9]: # Merge results and races on raceId and create a main dataframe
df = pd.merge(results, races[['raceId', 'year', 'circuitId']], on="raceId")

# create qualifying_time feature from q1, q2, q3
qualifying['qualifying_time'] = qualifying[['q1', 'q2', 'q3']].min(axis=1)

# adding qualifying_time feature to the main dataframe
df = pd.merge(df, qualifying[["raceId", "driverId", "constructorId",
    ↳ "qualifying_time"]], on=["raceId", "driverId", "constructorId"])

# fill the missing values if any with max value
df['qualifying_time'] = df.groupby('raceId')['qualifying_time'].
    ↳ transform(lambda x: x.fillna(x.max()))
```

```
[10]: # temporary data frame to generate grid_pos1, grid_pos2, ..., grid_pos10
    ↳ features
df_grid = pd.merge(races[['raceId', 'year', 'round']],
    ↳ results[['raceId', 'driverId', 'grid']] , on=['raceId'])
df_grid = df_grid.sort_values(by=['year', 'driverId', 'round'])

previous_rounds = 10

# generating and updating grid_pos1, grid_pos2, ..., grid_pos10 features
for i in range(1, previous_rounds + 1):
    df_grid[f'grid_pos{i}'] = df_grid.groupby(['year', 'driverId'])['grid'].
    ↳ shift(i)

df_grid = df_grid.reset_index(drop=True)

for i in range(1, previous_rounds + 1):
    df_grid.loc[df_grid['round'] <= i, f'grid_pos{i}'] = None

df_grid = df_grid.sort_values(by=['driverId', 'round'])

# Filling the None values with mean values
for i in range(1, previous_rounds + 1):
    if i == 1:
        df_grid[f'grid_pos{i}'].fillna(df_grid['grid'], inplace=True)
    else:
        df_grid[f'grid_pos{i}'].fillna(df_grid[[f'grid_pos{j}' for j in
    ↳ range(1, i)] + ['grid']].mean(axis=1), inplace=True)

# making grid_pos1, grid_pos2, ..., grid_pos10 features as int type
for i in range(1, previous_rounds + 1):
    df_grid[f'grid_pos{i}'] = df_grid[f'grid_pos{i}'].astype(int)
```

```
[11]: # temporary data frame to generate result_pos1, result_pos2, ..., result_pos10
      ↪ features
df_result = pd.merge(races[['raceId', 'year', 'round']],
      ↪ results[['raceId', 'driverId', 'position']] , on=['raceId'])
df_result = df_result.sort_values(by=['year', 'driverId', 'round'])

previous_rounds = 10

# generating and updating result_pos1, result_pos2, ..., result_pos10 features
for i in range(1, previous_rounds + 1):
    df_result[f'result_pos{i}'] = df_result.groupby(['year',
      ↪ 'driverId'])['position'].shift(i)

df_result = df_result.reset_index(drop=True)

for i in range(1, previous_rounds + 1):
    df_result.loc[df_result['round'] <= i, f'result_pos{i}'] = None

df_result = df_result.sort_values(by=['driverId', 'round'])

# Filling the None values with mean values
for i in range(1, previous_rounds + 1):
    if i == 1:
        df_result[f'result_pos{i}'].fillna(df_result['position'], inplace=True)
    else:
        df_result[f'result_pos{i}'].fillna(df_result[[f'result_pos{j}' for j in
      ↪ range(1, i)] + ['position']].mean(axis=1), inplace=True)

# making result_pos1, result_pos2, ..., result_pos10 features as int type
for i in range(1, previous_rounds + 1):
    df_result[f'result_pos{i}'] = df_result[f'result_pos{i}'].astype(int)
```

```
[12]: # adding grid_pos1, grid_pos2, ..., grid_pos10 feature to the main dataframe
df = pd.merge(df,
      ↪ df_grid[["raceId", "driverId", "grid_pos1", "grid_pos2", "grid_pos3", "grid_pos4", "grid_pos5", "g
      ↪ on=["raceId", "driverId"])

# adding result_pos1, result_pos2, ..., result_pos10 feature to the main
      ↪ dataframe
df = pd.merge(df,
      ↪ df_result[["raceId", "driverId", "result_pos1", "result_pos2", "result_pos3", "result_pos4", "res
      ↪ on=["raceId", "driverId"])

# position to points mapping
points_mapping = {
    1: 25, 2: 18, 3: 15, 4: 12, 5: 10, 6: 8, 7: 6, 8: 4, 9: 2, 10: 1
}
```

```
# modifying points according to the position using our mapping
df['points'] = df['position'].map(points_mapping).fillna(0)

# drop resultId, year, position as they are irrelevant in predictions
df = df.drop(['resultId', 'year', 'position'], axis=1)
```

1.4.4 4.4) Descriptive stats of the data and data exploration

```
[13]: df.head()
```

```
[13]:   raceId  driverId  constructorId  grid  points  laps  fastestLap  \
0      18         1           1      1    25.0    58         39
1      18         2           2      5    18.0    58         41
2      18         3           3      7    15.0    58         41
3      18         4           4     11    12.0    58         58
4      18         5           1      3    10.0    58         43

   fastestLapTime  fastestLapSpeed  circuitId  ...  result_pos1  result_pos2  \
0          87.452         218.300           1  ...           1           1
1          87.739         217.586           1  ...           2           2
2          88.090         216.719           1  ...           3           3
3          88.603         215.464           1  ...           4           4
4          87.418         218.385           1  ...           5           5

   result_pos3  result_pos4  result_pos5  result_pos6  result_pos7  \
0             1             1             1             1             1
1             2             2             2             2             2
2             3             3             3             3             3
3             4             4             4             4             4
4             5             5             5             5             5

   result_pos8  result_pos9  result_pos10
0             1             1             1
1             2             2             2
2             3             3             3
3             4             4             4
4             5             5             5

[5 rows x 31 columns]
```

```
[14]: df.dtypes
```

```
[14]: raceId          int64
      driverId       int64
      constructorId  int64
      grid           int64
```

```

points          float64
laps            int64
fastestLap      int64
fastestLapTime  float64
fastestLapSpeed float64
circuitId       int64
qualifying_time float64
grid_pos1       int64
grid_pos2       int64
grid_pos3       int64
grid_pos4       int64
grid_pos5       int64
grid_pos6       int64
grid_pos7       int64
grid_pos8       int64
grid_pos9       int64
grid_pos10      int64
result_pos1     int64
result_pos2     int64
result_pos3     int64
result_pos4     int64
result_pos5     int64
result_pos6     int64
result_pos7     int64
result_pos8     int64
result_pos9     int64
result_pos10    int64
dtype: object

```

```
[15]: df.describe()
```

```

[15]:
      count      raceId      driverId  constructorId      grid      points  \
count  8409.000000  8409.000000    8409.000000    8409.000000  8409.000000
mean    648.171483   355.634201     49.058628     10.911167    4.809609
std     427.310243   396.306270     75.318693      6.178142    7.086527
min       1.000000     1.000000      1.000000      0.000000    0.000000
25%     102.000000    14.000000      4.000000      6.000000    0.000000
50%     886.000000    37.000000      9.000000     11.000000    0.000000
75%     998.000000   822.000000     51.000000     16.000000    8.000000
max    1110.000000   858.000000    214.000000     24.000000   25.000000

      laps  fastestLap  fastestLapTime  fastestLapSpeed  circuitId  \
count  8409.000000  8409.000000    8409.000000    8409.000000  8409.000000
mean    52.434416   38.391247     670.015369     183.45365    19.355809
std     18.738098   20.325476    1756.901843      63.80533    20.863197
min       0.000000     0.000000      55.404000       0.00000     1.000000
25%     51.000000   24.000000      81.695000    188.01300     6.000000

```


50%	56.000000	43.000000	92.804000	202.41400	13.000000
75%	66.000000	53.000000	103.256000	214.81500	21.000000
max	87.000000	85.000000	6000.000000	257.32000	79.000000

	...	result_pos1	result_pos2	result_pos3	result_pos4	result_pos5	\
count	...	8409.000000	8409.000000	8409.000000	8409.000000	8409.000000	
mean	...	11.329766	11.338923	11.327268	11.309906	11.300987	
std	...	6.394765	6.347120	6.271143	6.186931	6.114017	
min	...	1.000000	1.000000	1.000000	1.000000	1.000000	
25%	...	6.000000	6.000000	6.000000	6.000000	6.000000	
50%	...	11.000000	11.000000	11.000000	11.000000	11.000000	
75%	...	17.000000	17.000000	17.000000	17.000000	16.000000	
max	...	24.000000	23.000000	23.000000	23.000000	23.000000	

	result_pos6	result_pos7	result_pos8	result_pos9	result_pos10
count	8409.000000	8409.000000	8409.000000	8409.000000	8409.000000
mean	11.287668	11.259841	11.256987	11.247711	11.224878
std	6.036886	5.943570	5.865866	5.772703	5.676413
min	1.000000	1.000000	1.000000	1.000000	1.000000
25%	6.000000	6.000000	6.000000	7.000000	7.000000
50%	11.000000	11.000000	11.000000	11.000000	11.000000
75%	16.000000	16.000000	16.000000	16.000000	16.000000
max	23.000000	23.000000	23.000000	23.000000	23.000000

[8 rows x 31 columns]

```
[16]: df.isna().sum()
```

```
[16]: raceId      0
      driverId    0
      constructorId  0
      grid        0
      points      0
      laps        0
      fastestLap   0
      fastestLapTime  0
      fastestLapSpeed  0
      circuitId    0
      qualifying_time  0
      grid_pos1    0
      grid_pos2    0
      grid_pos3    0
      grid_pos4    0
      grid_pos5    0
      grid_pos6    0
      grid_pos7    0
      grid_pos8    0
```

```

grid_pos9      0
grid_pos10     0
result_pos1    0
result_pos2    0
result_pos3    0
result_pos4    0
result_pos5    0
result_pos6    0
result_pos7    0
result_pos8    0
result_pos9    0
result_pos10   0
dtype: int64

```

```

[17]: print("Number of records = "+str(df.shape[0]))
      print("Number of features(including raceID) = "+str(df.shape[1]))

```

Number of records = 8409

Number of features(including raceID) = 31

```

[18]: # Count the number of races each driver has participated in
      race_counts = df['driverId'].value_counts()

      # Filter out drivers with less than 5 races
      eligible_drivers = race_counts[race_counts >= 10].index

      # Filter the main DataFrame for only eligible drivers
      df_eligible = df[df['driverId'].isin(eligible_drivers)]

      # Group by driverId and calculate the mean points for each eligible driver
      df_avg_points = df_eligible.groupby('driverId')['points'].mean().reset_index()

      # Sort the Series by average points in descending order to get the best drivers
      df_avg_points_sorted = df_avg_points.sort_values('points', ascending=False)

      # Select the top N drivers (adjust N based on your preference)
      top_n_drivers = 10
      df_top_n_avg_points = df_avg_points_sorted.head(top_n_drivers)

      # Merge with the original DataFrame to get additional information including
      #   driver names
      df_top_n_avg_info = pd.merge(df_top_n_avg_points, df_eligible, on='driverId',
      #   how='left')
      df_top_n_avg_info = pd.merge(df_top_n_avg_info, drivers[['driverId',
      #   'driverRef', 'forename', 'surname']], on='driverId', how='left')

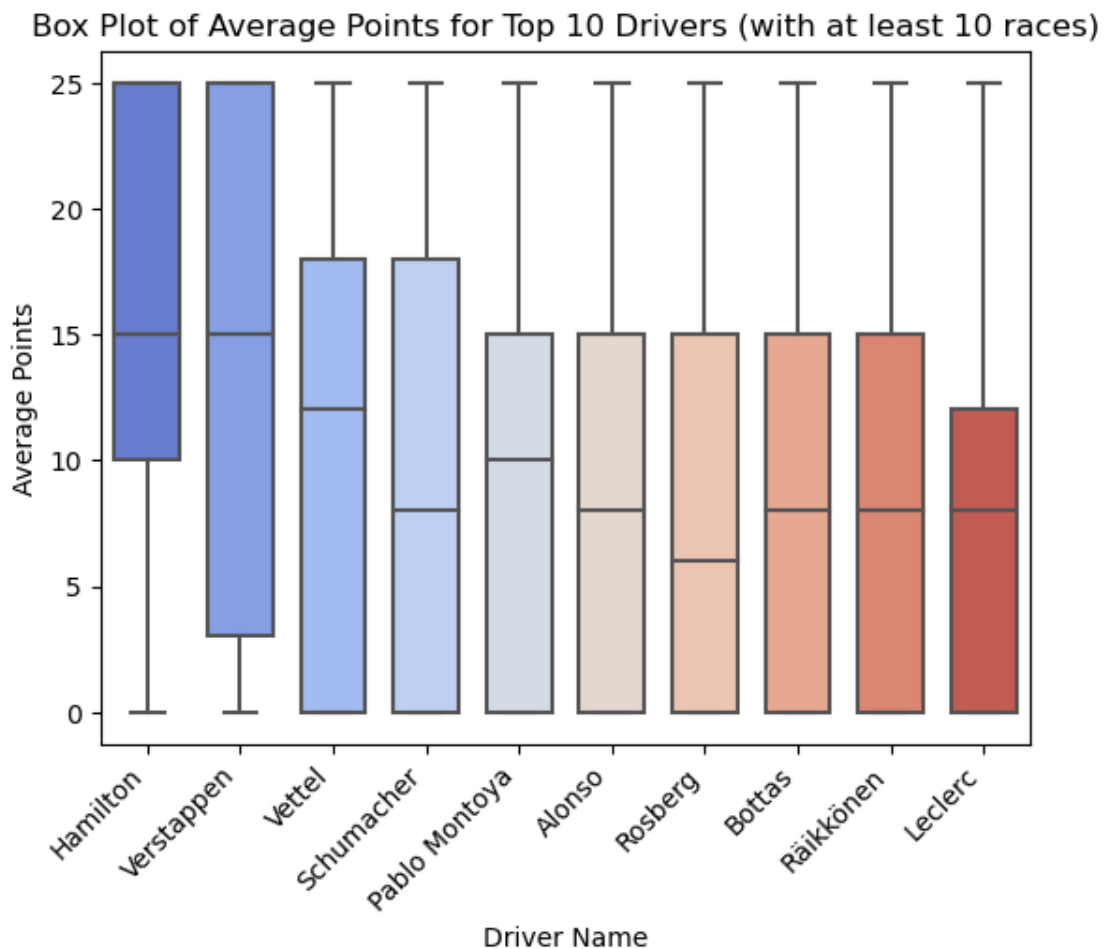
      # Create a box plot using Seaborn with switched axes and a cool color gradient

```

```

plt.figure()
sns.boxplot(x='surname', y='points_y', data=df_top_n_avg_info, width=0.7,
            palette="coolwarm")
plt.title('Box Plot of Average Points for Top {} Drivers (with at least 10
            races)'.format(top_n_drivers))
plt.xlabel('Driver Name')
plt.ylabel('Average Points')
plt.xticks(rotation=45, ha="right") # Rotate x-axis labels for better
            visibility
plt.show()

```



```

[19]: # Group by constructorId and calculate the mean points for each constructor
df_avg_points_constructor = df.groupby('constructorId')['points'].mean().
            reset_index()

```

```

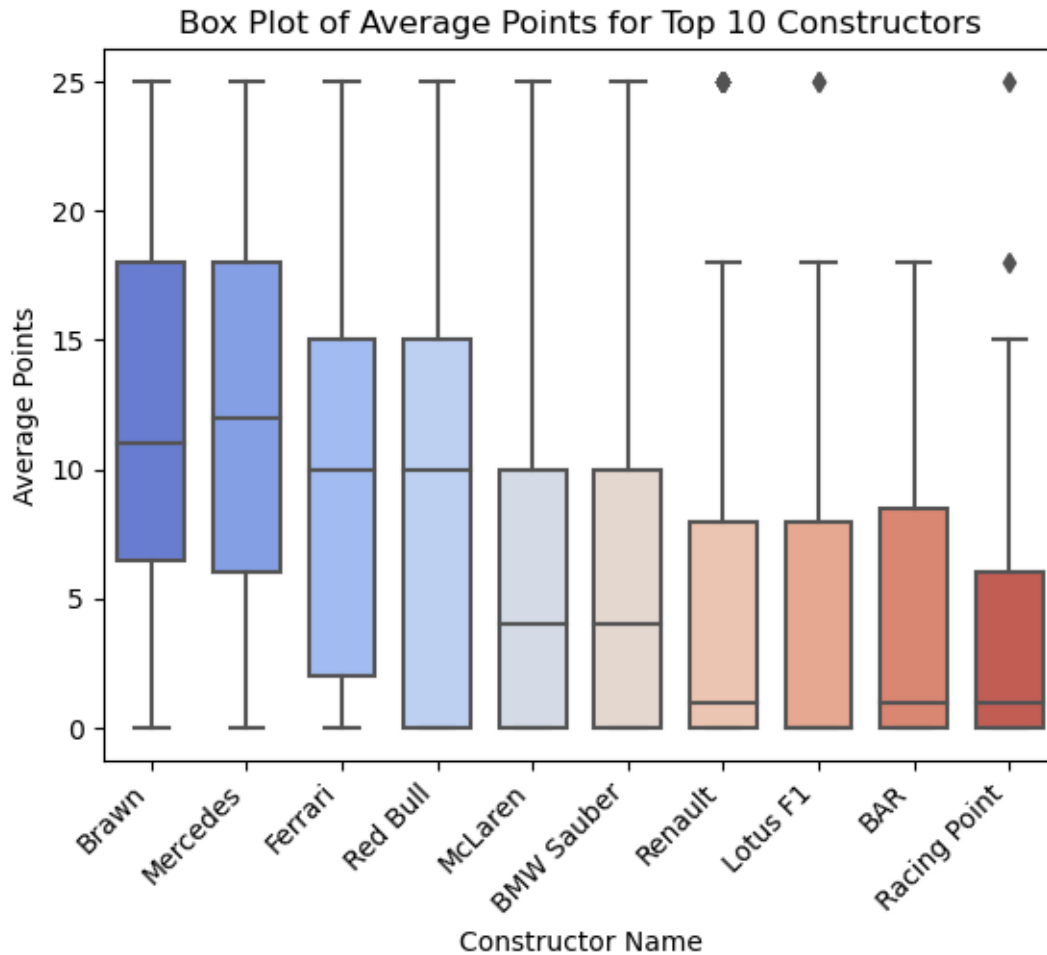
# Sort the DataFrame by average points in descending order to get the best
↳ constructors
df_avg_points_constructor_sorted = df_avg_points_constructor.
↳ sort_values('points', ascending=False)

# Select the top N constructors (adjust N based on your preference)
top_n_constructors = 10
df_top_n_avg_points_constructor = df_avg_points_constructor_sorted.
↳ head(top_n_constructors)

# Merge with the original DataFrame to get additional information including
↳ constructor names
df_top_n_avg_info_constructor = pd.merge(df_top_n_avg_points_constructor, df,
↳ on='constructorId', how='left')
df_top_n_avg_info_constructor = pd.merge(df_top_n_avg_info_constructor,
↳ constructors[['constructorId', 'name']], on='constructorId', how='left')

# Create a box plot using Seaborn with switched axes
plt.figure()
sns.boxplot(x='name', y='points_y', data=df_top_n_avg_info_constructor, width=0.
↳ 7, palette="coolwarm")
plt.title('Box Plot of Average Points for Top {} Constructors'.
↳ format(top_n_constructors))
plt.xlabel('Constructor Name')
plt.ylabel('Average Points')
plt.xticks(rotation=45, ha="right") # Rotate x-axis labels for better
↳ visibility
plt.show()

```



1.4.5 4.5) Feature Correlation Analysis


```
[20]: # Pair-wise Correlation of all the features
correlation = df.corr()

# Number of pairs of feaures which have correlation coefficient > 0.8
correlation_count = (np.sum(np.where(np.abs(correlation)>=0.8,1,0)) -
    ↳ correlation.shape[0])/2
print("Number of pairs of feaures which have correlation coefficient > 0.8 = "
    ↳ str(correlation_count))

# Correlation heat map plot
plt.title('Correlation plot')
print("Correlation plot:")
sns.heatmap(correlation);
```

Number of pairs of feaures which have correlation coefficient > 0.8 = 2.0

Correlation plot:



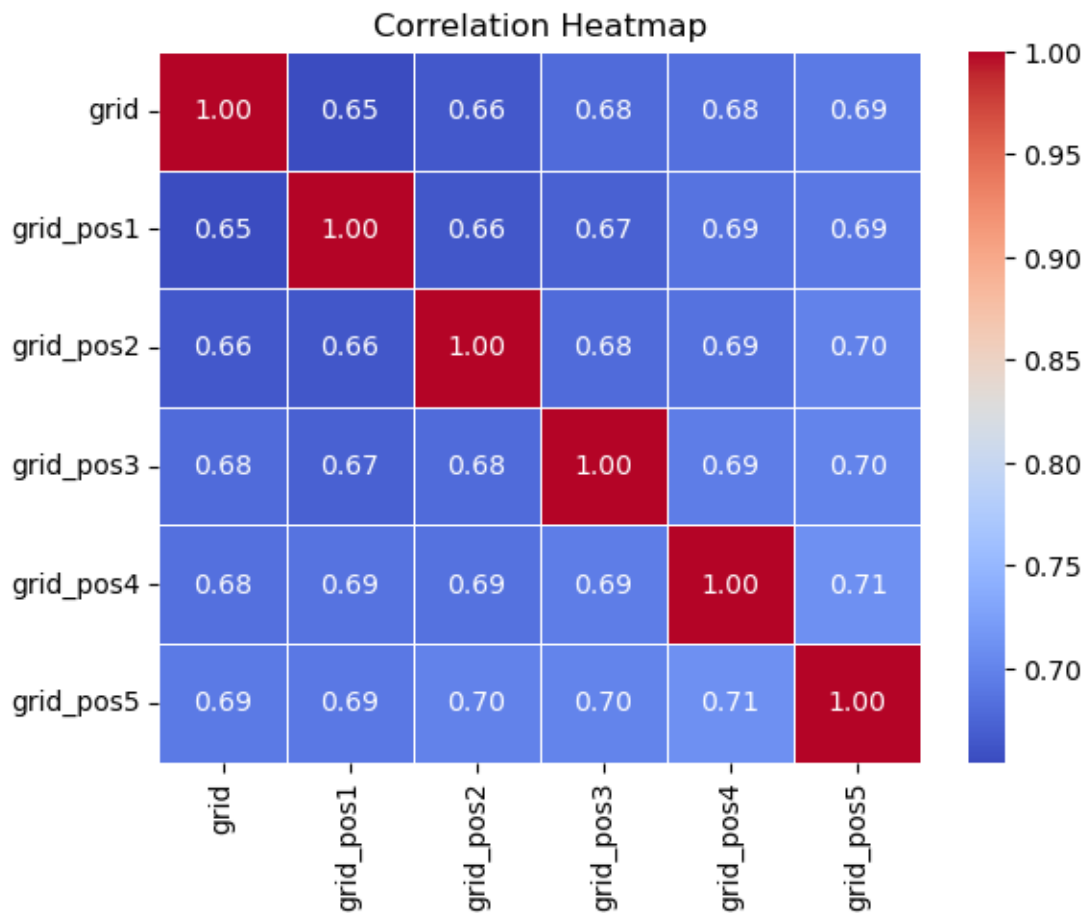
We can observe that the the features are mostly independent and not correlated with eachother. So data reduction using pca or svd is not necessary

```
[21]: # correlation between 'grid_pos1', 'grid_pos2', 'grid_pos3', 'grid_pos4',
      ↪ 'grid_pos5' and 'grid'
columns_of_interest = ['grid', 'grid_pos1', 'grid_pos2', 'grid_pos3',
      ↪ 'grid_pos4', 'grid_pos5']
df_selected = df[columns_of_interest]

# Calculate the correlation matrix
correlation_matrix = df_selected.corr()

# Create a heatmap using Seaborn
plt.figure()
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt='.2f',
      ↪ linewidths=.5)
```

```
plt.title('Correlation Heatmap')
plt.show()
```




```
[22]: # correlation between 'result_pos1', 'result_pos2', 'result_pos3',
      ↪ 'result_pos4', 'result_pos5' and 'points'
columns_of_interest = ['points', 'result_pos1', 'result_pos2', 'result_pos3',
      ↪ 'result_pos4', 'result_pos5']
df_selected = df[columns_of_interest]

# Calculate the correlation matrix
correlation_matrix = df_selected.corr()

# Create a heatmap using Seaborn
plt.figure()
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt='.2f',
      ↪ linewidths=.5)
plt.title('Correlation Heatmap')
```

```
plt.show()
```



1.4.6 4.6) Encoding categorical features

```
[23]: # OneHot encoding 'driverId', 'constructorId', 'circuitId' (Qualitative Nominal) ↵  
      ↪ features  
enc = OneHotEncoder()  
encoded = enc.fit_transform(df[['driverId', 'constructorId', 'circuitId']])  
df_enc = df.drop(['driverId', 'constructorId', 'circuitId'], axis=1)  
df_enc = pd.concat([df_enc, pd.DataFrame(encoded.toarray())], axis=1)  
df_enc.columns = df_enc.columns.astype(str)
```


1.4.7 4.7) Data Split to train and test

```
[24]: # get unique raceIds
unique_raceIds = df_enc['raceId'].unique()

# Split the unique raceIds to train(80%) and test(20%)
train_raceIds, test_raceIds = train_test_split(unique_raceIds, test_size=0.2,
random_state=42)

# Split the recored corresponding to train_raceIds and test_raceIds
train_set = df_enc[df_enc['raceId'].isin(train_raceIds)]
test_set = df_enc[df_enc['raceId'].isin(test_raceIds)]

train_raceId_list = df_enc[df_enc['raceId'].isin(train_raceIds)]['raceId']
test_raceId_list = df_enc[df_enc['raceId'].isin(test_raceIds)]['raceId']

# Split as X(features) and Y(target)
X_train = train_set.drop(columns=['raceId', 'points'])
Y_train = train_set['points']

# Split as X(features) and Y(target)
X_test = test_set.drop(columns=['raceId', 'points'])
Y_test = test_set['points']
```

```
[25]: print("Train data samples: " +str(X_train.shape[0]))
print("Test data samples: " +str(X_test.shape[0]))
```

Train data samples: 6720

Test data samples: 1689

1.5 5) Models

1.5.1 5.1) Linear Regression

```
[26]: # Linear regression
linear_reg = LinearRegression().fit(X_train, Y_train)

# Train predictions
Y_train_pred = linear_reg.predict(X_train)
Y_train_pred = attach_raceId_to_predictions(train_raceId_list, Y_train_pred)

# Test predictions
Y_test_pred = linear_reg.predict(X_test)
Y_test_pred = attach_raceId_to_predictions(test_raceId_list, Y_test_pred)

# Calculate R2 score for train and test
linear_reg_train_score = r2_score(Y_train, Y_train_pred)
linear_reg_test_score = r2_score(Y_test, Y_test_pred)
```

```

# Calculate RootMeanSquareError for train and test
linear_reg_train_rmse = np.sqrt(mean_squared_error(Y_train, Y_train_pred))
linear_reg_test_rmse = np.sqrt(mean_squared_error(Y_test, Y_test_pred))

# Print the evaluation metics
print("Train R2 score = "+str(linear_reg_train_score))
print("Test R2 score = "+str(linear_reg_test_score))
print("Train RootMeanSquareError = "+str(linear_reg_train_rmse))
print("Test RootMeanSquareError = "+str(linear_reg_test_rmse))

```

```

Train R2 score = 0.5642283556406484
Test R2 score = 0.6269702783572284
Train RootMeanSquareError = 4.676005748703134
Test RootMeanSquareError = 4.334312409318795

```

1.5.2 5.2) Ridge Regression

```

[27]: # Ridge regression
ridge_reg = Ridge(alpha=0.001).fit(X_train, Y_train)

# Train predictions
Y_train_pred = ridge_reg.predict(X_train)
Y_train_pred = attach_raceId_to_predictions(train_raceId_list, Y_train_pred)

# Test predictions
Y_test_pred = ridge_reg.predict(X_test)
Y_test_pred = attach_raceId_to_predictions(test_raceId_list, Y_test_pred)

# Calculate R2 score for train and test
ridge_reg_train_score = r2_score(Y_train, Y_train_pred)
ridge_reg_test_score = r2_score(Y_test, Y_test_pred)

# Calculate RootMeanSquareError for train and test
ridge_reg_train_rmse = np.sqrt(mean_squared_error(Y_train, Y_train_pred))
ridge_reg_test_rmse = np.sqrt(mean_squared_error(Y_test, Y_test_pred))

# Print the evaluation metics
print("Train R2 score = "+str(ridge_reg_train_score))
print("Test R2 score = "+str(ridge_reg_test_score))
print("Train RootMeanSquareError = "+str(ridge_reg_train_rmse))
print("Test RootMeanSquareError = "+str(ridge_reg_test_rmse))

```

```

Train R2 score = 0.5642283556406484
Test R2 score = 0.6269702783572284
Train RootMeanSquareError = 4.676005748703134
Test RootMeanSquareError = 4.334312409318795

```

1.5.3 5.3) Lasso Regression

```
[28]: # Lasso regression
lasso_reg = Lasso(alpha=0.001).fit(X_train, Y_train)

# Train predictions
Y_train_pred = lasso_reg.predict(X_train)
Y_train_pred = attach_raceId_to_predictions(train_raceId_list, Y_train_pred)

# Test predictions
Y_test_pred = lasso_reg.predict(X_test)
Y_test_pred = attach_raceId_to_predictions(test_raceId_list, Y_test_pred)

# Calculate R2 score for train and test
lasso_reg_train_score = r2_score(Y_train, Y_train_pred)
lasso_reg_test_score = r2_score(Y_test, Y_test_pred)

# Calculate RootMeanSquareError for train and test
lasso_reg_train_rmse = np.sqrt(mean_squared_error(Y_train, Y_train_pred))
lasso_reg_test_rmse = np.sqrt(mean_squared_error(Y_test, Y_test_pred))

# Print the evaluation metics
print("Train R2 score = "+str(lasso_reg_train_score))
print("Test R2 score = "+str(lasso_reg_test_score))
print("Train RootMeanSquareError = "+str(lasso_reg_train_rmse))
print("Test RootMeanSquareError = "+str(lasso_reg_test_rmse))
```

Train R2 score = 0.5620751959433383

Test R2 score = 0.6258651792764083

Train RootMeanSquareError = 4.687543650590411

Test RootMeanSquareError = 4.3407278531196845

1.5.4 5.4) Decision Tree

```
[29]: # hyper-parameter tuning on max_depth parameter
param_grid = {
    'max_depth': [2, 4, 6, 8, 10, 12, 14, 16, 32, 64],
}

# Decision tree regression
decisiontree_reg = DecisionTreeRegressor(random_state=42)

# Grid search on max_depth hyper-parameter tuning
decisiontree_reg_grid_search = GridSearchCV(estimator= decisiontree_reg,
    ↪ param_grid=param_grid, scoring='r2')
decisiontree_reg_grid_result = decisiontree_reg_grid_search.fit(X_train,
    ↪ Y_train)
```

```

# Get best model and it's hyper paramets
decisiontree_reg_best_params = decisiontree_reg_grid_result.best_params_
decisiontree_reg_best_model = decisiontree_reg_grid_result.best_estimator_

# Train predictions on best model
Y_train_pred = decisiontree_reg_best_model.predict(X_train)
Y_train_pred = attach_raceId_to_predictions(train_raceId_list, Y_train_pred)

# Test predictions on best model
Y_test_pred = decisiontree_reg_best_model.predict(X_test)
Y_test_pred = attach_raceId_to_predictions(test_raceId_list, Y_test_pred)

# Calculate R2 score for train and test
decisiontree_reg_train_score = r2_score(Y_train, Y_train_pred)
decisiontree_reg_test_score = r2_score(Y_test, Y_test_pred)

# Calculate RootMeanSquareError for train and test
decisiontree_reg_train_rmse = np.sqrt(mean_squared_error(Y_train, Y_train_pred))
decisiontree_reg_test_rmse = np.sqrt(mean_squared_error(Y_test, Y_test_pred))

# Print the evaluation metics
print("Train R2 score = "+str(decisiontree_reg_train_score))
print("Test R2 score = "+str(decisiontree_reg_test_score))
print("Train RootMeanSquareError = "+str(decisiontree_reg_train_rmse))
print("Test RootMeanSquareError = "+str(decisiontree_reg_test_rmse))
print("Best Prameters: "+str(decisiontree_reg_best_params))

```

```

Train R2 score = 0.6671885375302121
Test R2 score = 0.615166879664214
Train RootMeanSquareError = 4.086435896148786
Test RootMeanSquareError = 4.402351489630001
Best Prameters: {'max_depth': 6}

```

1.5.5 5.5) Random Forest

```

[30]: # hyper-parameter tuning on max_depth and n_estimators parameter
param_grid = {
    'n_estimators': [50, 100, 150, 200, 250],
    'max_depth': [2, 4, 6, 8, 10, 12, 14, 16, 32, 64],
}

# Random Forest Regressor
rf_model = RandomForestRegressor(random_state=42)

# Grid search on max_depth hyper-parameter tuning
rf_model_grid_search = GridSearchCV(estimator=rf_model, param_grid=param_grid,
    scoring='r2', cv=5)

```

```

rf_model_grid_result = rf_model_grid_search.fit(X_train, Y_train)

# Get best model and it's hyper params
rf_model_best_params = rf_model_grid_result.best_params_
rf_model_best_model = rf_model_grid_result.best_estimator_

# Train predictions on best model
Y_train_pred = rf_model_best_model.predict(X_train)
Y_train_pred = attach_raceId_to_predictions(train_raceId_list, Y_train_pred)

# Test predictions on best model
Y_test_pred = rf_model_best_model.predict(X_test)
Y_test_pred = attach_raceId_to_predictions(test_raceId_list, Y_test_pred)

# Calculate R2 score for train and test
rf_model_train_score = r2_score(Y_train, Y_train_pred)
rf_model_test_score = r2_score(Y_test, Y_test_pred)

# Calculate RootMeanSquareError for train and test
rf_model_train_rmse = np.sqrt(mean_squared_error(Y_train, Y_train_pred))
rf_model_test_rmse = np.sqrt(mean_squared_error(Y_test, Y_test_pred))

# Print the evaluation metics
print("Train R2 score = "+str(rf_model_train_score))
print("Test R2 score = "+str(rf_model_test_score))
print("Train RootMeanSquareError = "+str(rf_model_train_rmse))
print("Test RootMeanSquareError = "+str(rf_model_test_rmse))
print("Best Prameters: "+str(rf_model_best_params))

```

Train R2 score = 0.9596860884772198
 Test R2 score = 0.6946634752440572
 Train RootMeanSquareError = 1.422240435771272
 Test RootMeanSquareError = 3.9213704543665497
 Best Prameters: {'max_depth': 16, 'n_estimators': 250}

1.5.6 5.6) XGBoost

```

[31]: # hyper-parameter tuning on these parameter
param_grid = {
    'max_depth': [4, 6],
    'learning_rate': [0.1, 0.01],
    'n_estimators': [50, 100, 150],
    'min_child_weight': [1, 2, 3],
    'subsample': [0.9],
    'colsample_bytree': [0.9],
}

```

```

# XGBoost Regressor
xgb_regressor = xgb.XGBRegressor(objective='reg:squarederror')

# 5-fold cross validation and Grid search on max_depth hyper-parameter tuning
kf = KFold(n_splits=5, shuffle=True, random_state=42)
xgb_regressor_grid_search = GridSearchCV(estimator=xgb_regressor,
    param_grid=param_grid, scoring='neg_mean_squared_error', cv=kf)
xgb_regressor_grid_result = xgb_regressor_grid_search.fit(X_train, Y_train)

# Get best model and it's hyper params
xgb_regressor_best_params = xgb_regressor_grid_result.best_params_
xgb_regressor_best_model = xgb_regressor_grid_result.best_estimator_

# Train predictions on best model
Y_train_pred = xgb_regressor_best_model.predict(X_train)
Y_train_pred = attach_raceId_to_predictions(train_raceId_list, Y_train_pred)

# Test predictions on best model
Y_test_pred = xgb_regressor_best_model.predict(X_test)
Y_test_pred = attach_raceId_to_predictions(test_raceId_list, Y_test_pred)

# Calculate R2 score for train and test
xgb_regressor_train_score = r2_score(Y_train, Y_train_pred)
xgb_regressor_test_score = r2_score(Y_test, Y_test_pred)

# Calculate RootMeanSquareError for train and test
xgb_regressor_train_rmse = np.sqrt(mean_squared_error(Y_train, Y_train_pred))
xgb_regressor_test_rmse = np.sqrt(mean_squared_error(Y_test, Y_test_pred))

# Print the evaluation metics
print("Train R2 score = "+str(xgb_regressor_train_score))
print("Test R2 score = "+str(xgb_regressor_test_score))
print("Train RootMeanSquareError = "+str(xgb_regressor_train_rmse))
print("Test RootMeanSquareError = "+str(xgb_regressor_test_rmse))
print("Best Prameters: "+str(xgb_regressor_best_params))

```

Train R2 score = 0.8991484910371411

Test R2 score = 0.7305909517439093

Train RootMeanSquareError = 2.24950391356469

Test RootMeanSquareError = 3.6834483290333226


Best Prameters: {'colsample_bytree': 0.9, 'learning_rate': 0.1, 'max_depth': 6, 'min_child_weight': 3, 'n_estimators': 100, 'subsample': 0.9}

1.6 6) Results, Models Evaluations and Conclusion

1.6.1 6.1) R2 score comparion plot

```
[32]: # different model comparisons with R2 metrics
models = ['Linear Regression', 'Ridge Regression', 'Lasso Regression',
          ↪ 'Decision Tree', 'Random Forest', 'XGBoost']
train_r2 = [linear_reg_train_score, ridge_reg_train_score,
          ↪ lasso_reg_train_score, decisiontree_reg_train_score, rf_model_train_score,
          ↪ xgb_regressor_train_score]
test_r2 = [linear_reg_test_score, ridge_reg_test_score, lasso_reg_test_score,
          ↪ decisiontree_reg_test_score, rf_model_test_score, xgb_regressor_test_score]

plt.plot(models, train_r2, label = 'train r2 score')
plt.plot(models, test_r2, label = 'test r2 score')
plt.title('R2 score')
plt.xlabel('Model Name')
plt.ylabel('R2 Score')
plt.xticks(rotation=45, ha="right") # Rotate x-axis labels for better
          ↪ visibility
plt.legend()
plt.show()
```




1.6.2 6.2) RMSE comparion plot

```
[33]: # different model comparisions MSE metrics
models = ['Linear Regression', 'Ridge Regression', 'Lasso Regression',
          'Decision Tree', 'Random Forest', 'XGBoost']
train_r2 = [linear_reg_train_rmse, ridge_reg_train_rmse, lasso_reg_train_rmse,
            decisiontree_reg_train_rmse, rf_model_train_rmse, xgb_regressor_train_rmse]
test_r2 = [linear_reg_test_rmse, ridge_reg_test_rmse, lasso_reg_test_rmse,
            decisiontree_reg_test_rmse, rf_model_test_rmse, xgb_regressor_test_rmse]

plt.plot(models, train_r2, label = 'train mse')
plt.plot(models, test_r2, label = 'test mse')
plt.title('Root Meam Square Error')
plt.xlabel('Model Name')
plt.ylabel('Root Meam Square Error')
```



```
plt.xticks(rotation=45, ha="right") # Rotate x-axis labels for better
    ↪ visibility
plt.legend()
plt.show()
```




1.6.3 6.3) Decision tree top 5 important features

```
[34]: # Get feature importances
decisiontree_reg_importances = decisiontree_reg_best_model.feature_importances_
indices = np.argsort(decisiontree_reg_importances)[::-1]

# Get top 5 feature indices and importances
top_5_indices = indices[:5]
top_5_importances = decisiontree_reg_importances[top_5_indices]
top_5_feature_names = [X_train.columns[i] for i in top_5_indices]
```

```
# Plot the top 5 feature importances
plt.figure()
plt.title("Top 5 Feature Importances - Decision Tree")
plt.bar(range(len(top_5_importances)), top_5_importances, align="center")
plt.xticks(range(len(top_5_importances)), top_5_feature_names, rotation=45)
plt.xlabel("Features")
plt.ylabel("Importance")
plt.tight_layout()
plt.show()
```



1.6.4 6.4) Random Forest top 5 important features

```
[35]: # Get feature importances
rf_model_importances = rf_model_best_model.feature_importances_
indices = np.argsort(rf_model_importances)[::-1]


# Get top 5 feature indices and importances
top_5_indices = indices[:5]
top_5_importances = rf_model_importances[top_5_indices]
```

```

top_5_feature_names = [X_train.columns[i] for i in top_5_indices]

# Plot the top 5 feature importances
plt.figure()
plt.title("Top 5 Feature Importances - Random Forest")
plt.bar(range(len(top_5_importances)), top_5_importances, align="center")
plt.xticks(range(len(top_5_importances)), top_5_feature_names, rotation=45)
plt.xlabel("Features")
plt.ylabel("Importance")
plt.tight_layout()
plt.show()

```



1.6.5 6.3) XGBoost top 5 important features

```

[36]: # Get feature importances
xgb_regressor_importances = xgb_regressor_best_model.feature_importances_
indices = np.argsort(xgb_regressor_importances)[::-1]

# Get top 5 feature indices and importances
top_5_indices = indices[:5]


```

```

top_5_importances = xgb_regressor_importances[top_5_indices]
top_5_feature_names = [X_train.columns[i] for i in top_5_indices]

# Plot the top 5 feature importances
plt.figure()
plt.title("Top 5 Feature Importances - XG Boost")
plt.bar(range(len(top_5_importances)), top_5_importances, align="center")
plt.xticks(range(len(top_5_importances)), top_5_feature_names, rotation=45)
plt.xlabel("Features")
plt.ylabel("Importance")
plt.tight_layout()
plt.show()

```



1.6.6 6.4) Final remarks and conclusion

1. Since it's a regression problem, we wanted to start with basic regression model and go to more complex models. That's why we trained our data on 'Linear Regression', 'Ridge Regression', 'Lasso Regression', 'Decision Tree', 'Random Forest' and 'XGBoost' until we got the reasonably performing model
2. Since it's a regression problem we choose RMSE and R2 score as metrics for model evaluations
3. Based on the R2 scores and RMSE values we can see that 'XGBoost' outperforms all the

other models 'Linear Regression', 'Ridge Regression', 'Lasso Regression', 'Decision Tree', 'Random Forest'

4. 'Linear Regression', 'Ridge Regression', 'Lasso Regression' didn't give good results the R2 score are less than 0.6 and have high RMSE
5. There is a slight improvement when we come to 'Decision Tree' because of increase in model complexity
6. Performance of ensemble methods 'Random Forest' and 'XGBoost' are higher than 'Decision Tree' which is expected and between 'Random Forest' and 'XGBoost', 'XGBoost' outperformed 'Random Forest'
7. Given the complexity of the problem we are trying to study these results makes sense because as the model complexity increases we observe good results
8. From the feature importance of all 'Decision Tree', 'Random Forest' and 'XGBoost' models we can see that features grid, laps, result_pos10, result_pos9, result_pos8 seems to have high significance while deciding the points in the race
9. This makes sense because all grid, laps, result_pos10, result_pos9, result_pos8 in some way indicates the performance of the race car and driver